

Assignment 3: Language Models: Auto-Complete

In this assignment, you will build an auto-complete system. Auto-complete system is something you may see every day

- When you google something, you often have suggestions to help you complete your search.
- When you are writing an email, you get suggestions telling you possible endings to your sentence.

By the end of this assignment, you will develop a prototype of such a system.



stanford is be

stanford is **better than** harvard

stanford is **best known for**

is stanford **better than** ivy league

is stanford **better than** berkeley

Important Note on Submission to the AutoGrader

Before submitting your assignment to the AutoGrader, please make sure you are not doing the following:

1. You have not added any *extra* `print` statement(s) in the assignment.
2. You have not added any *extra* code cell(s) in the assignment.
3. You have not changed any of the function parameters.
4. You are not using any global variables inside your graded exercises. Unless specifically instructed to do so, please refrain from it and use the local variables instead.
5. You are not changing the assignment code where it is not required, like creating *extra* variables.

If you do any of the following, you will get something like, Grader Error: Grader feedback not found (or similarly unexpected) error upon submitting your assignment. Before asking for help/debugging the errors in your assignment, check for these first. If this is the case, and you don't remember the changes you have made, you can get a fresh copy of the assignment by following these [instructions \(https://www.coursera.org/learn/probabilistic-models-in-nlp/supplement/saGQf/how-to-refresh-your-workspace\)](https://www.coursera.org/learn/probabilistic-models-in-nlp/supplement/saGQf/how-to-refresh-your-workspace).

Table of Contents

- [1 - Load and Preprocess Data](#)
 - [1.1 - Load the Data](#)
 - [1.2 - Pre-process the Data](#)
 - [Exercise 1- split to sentences \(UNQ C1\)](#)
 - [Exercise 2 - tokenize sentences \(UNQ C2\)](#)
 - [Exercise 3 - get tokenized data \(UNQ C3\)](#)
 - [Exercise 4 - count words \(UNQ C4\)](#)
 - [Exercise 5 - get words with nplus frequency \(UNQ C5\)](#)
 - [Exercise 6 - replace oov words by unk \(UNQ C6\)](#)
 - [Exercise 7 - preprocess data \(UNQ C7\)](#)
- [2 - Develop n-gram based Language Models](#)
 - [Exercise 8 - count n_grams \(UNQ C8\)](#)
 - [Exercise 9 - estimate probability \(UNQ C9\)](#)
- [3 - Perplexity](#)
 - [Exercise 10 - calculate perplexity \(UNQ C10\)](#)
- [4 - Build an Auto-complete System](#)
 - [Exercise 11 - suggest a word \(UNQ C11\)](#)

A key building block for an auto-complete system is a language model. A language model assigns the probability to a sequence of words, in a way that more "likely" sequences receive higher scores. For example,

"I have a pen" is expected to have a higher probability than "I am a pen" since the first one seems to be a more natural sentence in the real world.

You can take advantage of this probability calculation to develop an auto-complete system. Suppose the user typed

"I eat scrambled" Then you can find a word x such that "I eat scrambled x " receives the highest probability. If x = "eggs", the sentence would be "I eat scrambled eggs"

While a variety of language models have been developed, this assignment uses **N-grams**, a simple but powerful method for language modeling.

- N-grams are also used in machine translation and speech recognition.

Here are the steps of this assignment:

1. Load and preprocess data
 - Load and tokenize data.
 - Split the sentences into train and test sets.
 - Replace words with a low frequency by an unknown marker `<unk>`.
2. Develop N-gram based language models
 - Compute the count of n-grams from a given data set.
 - Estimate the conditional probability of a next word with k-smoothing.
3. Evaluate the N-gram models by computing the perplexity score.

```
In [1]: import math
import random
import numpy as np
import pandas as pd
import nltk
nltk.download('punkt')

import w3_unittest
nltk.data.path.append('.')
```

```
[nltk_data] Downloading package punkt to /home/jovyan/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

1 - Load and Preprocess Data

1.1 - Load the Data

You will use twitter data. Load the data and view the first few sentences by running the next cell.

Notice that data is a long string that contains many many tweets. Observe that there is a line break "\n" between tweets.

```
In [2]: with open("../data/en_US.twitter.txt", "r") as f:
        data = f.read()
        print("Data type:", type(data))
        print("Number of letters:", len(data))
        print("First 300 letters of the data")
        print("-----")
        display(data[0:300])
        print("-----")

        print("Last 300 letters of the data")
        print("-----")
        display(data[-300:])
        print("-----")
```

```
Data type: <class 'str'>
Number of letters: 3335477
First 300 letters of the data
-----
```

```
"How are you? Btw thanks for the RT. You gonna be in DC anytime soon? Love to see you. Been way, way too long.\nWhen you meet someone special... you'll know. Your heart will beat more rapidly and you'll smile for no reason.\nthe y've decided its more fun if I don't.\nSo Tired D; Played Lazer Tag & Ran A"
```

```
-----
Last 300 letters of the data
-----
```

```
"ust had one a few weeks back....hopefully we will be back soon! wish you the best yo\nColombia is with an 'o'..."": We now ship to 4 countries in South America (fist pump). Please welcome Columbia to the Stunner Family"\n#GutsiestMovesYouCanMake Giving a cat a bath.\nCoffee after 5 was a TERRIBLE idea.\n"
```

```
-----
```

1.2 - Pre-process the Data

Preprocess this data with the following steps:

1. Split data into sentences using "\n" as the delimiter.
2. Split each sentence into tokens. Note that in this assignment we use "token" and "words" interchangeably.
3. Assign sentences into train or test sets.
4. Find tokens that appear at least N times in the training data.
5. Replace tokens that appear less than N times by <unk>

Note: we omit validation data in this exercise.

- In real applications, we should hold a part of data as a validation set and use it to tune our training.
- We skip this process for simplicity.

Exercise 1- split_to_sentences

Hints

```
In [3]: # UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
### GRADED_FUNCTION: split_to_sentences ###
def split_to_sentences(data):
    """
    Split data by linebreak "\n"

    Args:
        data: str

    Returns:
        A list of sentences
    """
    ### START CODE HERE (Replace instances of 'None' with your code) ###
    sentences = data.split('\n')
    ### END CODE HERE ###

    # Additional cleaning (This part is already implemented)
    # - Remove leading and trailing spaces from each sentence
    # - Drop sentences if they are empty strings.
    sentences = [s.strip() for s in sentences]
    sentences = [s for s in sentences if len(s) > 0]

    return sentences
```

```
In [4]: # test your code
x = """
I have a pen.\nI have an apple. \nAh\nApple pen.\n
"""
print(x)

split_to_sentences(x)
```

I have a pen.
I have an apple.
Ah
Apple pen.

```
Out[4]: ['I have a pen.', 'I have an apple.', 'Ah', 'Apple pen.']
```

Expected answer:

```
['I have a pen.', 'I have an apple.', 'Ah', 'Apple pen.']
```

```
In [5]: # Test your function
w3_unittest.test_split_to_sentences(split_to_sentences)
```

All tests passed

Exercise 2 - tokenize_sentences

The next step is to tokenize sentences (split a sentence into a list of words).

- Convert all tokens into lower case so that words which are capitalized (for example, at the start of a sentence) in the original text are treated the same as the lowercase versions of the words.
- Append each tokenized list of words into a list of tokenized sentences.

Hints

```
In [6]: import re

# UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
### GRADED_FUNCTION: tokenize_sentences ###
def tokenize_sentences(sentences):
    """
    Tokenize sentences into tokens (words)

    Args:
        sentences: List of strings

    Returns:
        List of lists of tokens
    """

    # Initialize the List of Lists of tokenized sentences
    tokenized_sentences = []
    ### START CODE HERE (Replace instances of 'None' with your code) ###

    # Go through each sentence
    for sentence in sentences:

        # Convert to Lowercase Letters
        sentence = sentence.lower()

        # Convert into a List of words
        tokenized = nltk.word_tokenize(sentence)

        # append the List of words to the List of Lists
        tokenized_sentences.append(tokenized)

    ### END CODE HERE ###

    return tokenized_sentences
```

```
In [7]: # test your code
sentences = ["Sky is blue.", "Leaves are green.", "Roses are red."]
tokenize_sentences(sentences)
```

```
Out[7]: [['sky', 'is', 'blue', '.'],
          ['leaves', 'are', 'green', '.'],
          ['roses', 'are', 'red', '.']]
```

Expected output

```
[['sky', 'is', 'blue', '.'],  
 ['leaves', 'are', 'green', '.'],  
 ['roses', 'are', 'red', '.']]
```

```
In [8]: # Test your function  
w3_unittest.test_tokenize_sentences(tokenize_sentences)
```

All tests passed

Exercise 3 - get_tokenized_data

Use the two functions that you have just implemented to get the tokenized data.

- split the data into sentences
- tokenize those sentences

```
In [10]: # UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)  
### GRADED_FUNCTION: get_tokenized_data ###  
def get_tokenized_data(data):  
    """  
    Make a list of tokenized sentences  
  
    Args:  
        data: String  
  
    Returns:  
        List of lists of tokens  
    """  
    ### START CODE HERE (Replace instances of 'None' with your code) ###  
  
    # Get the sentences by splitting up the data  
    sentences = split_to_sentences(data)  
  
    # Get the list of lists of tokens by tokenizing the sentences  
    tokenized_sentences = tokenize_sentences(sentences)  
  
    ### END CODE HERE ###  
  
    return tokenized_sentences
```

```
In [11]: # test your function  
x = "Sky is blue.\nLeaves are green\nRoses are red."  
get_tokenized_data(x)
```

```
Out[11]: [['sky', 'is', 'blue', '.'],  
 ['leaves', 'are', 'green'],  
 ['roses', 'are', 'red', '.']]
```

Expected outcome

```
[['sky', 'is', 'blue', '.'],  
 ['leaves', 'are', 'green'],  
 ['roses', 'are', 'red', '.']]
```

```
In [12]: # Test your function  
w3_unittest.test_get_tokenized_data(get_tokenized_data)
```

All tests passed

Split into train and test sets

Now run the cell below to split data into training and test sets.

```
In [13]: tokenized_data = get_tokenized_data(data)  
random.seed(87)  
random.shuffle(tokenized_data)  
  
train_size = int(len(tokenized_data) * 0.8)  
train_data = tokenized_data[0:train_size]  
test_data = tokenized_data[train_size:]
```

```
In [14]: print("{} data are split into {} train and {} test set".format(  
        len(tokenized_data), len(train_data), len(test_data)))  
  
print("First training sample:")  
print(train_data[0])  
  
print("First test sample")  
print(test_data[0])
```

47961 data are split into 38368 train and 9593 test set

First training sample:

```
['i', 'personally', 'would', 'like', 'as', 'our', 'official', 'glove', 'of',  
 'the', 'team', 'local', 'company', 'and', 'quality', 'production']
```

First test sample

```
['that', 'picture', 'i', 'just', 'seen', 'whoa', 'dere', '!', '!', '>', '>',  
 '>', '>', '>', '>', '>']
```

Expected output

47961 data are split into 38368 train and 9593 test set

First training sample:

```
['i', 'personally', 'would', 'like', 'as', 'our', 'official', 'glove',  
 'of', 'the', 'team', 'local', 'company', 'and', 'quality', 'production']
```

First test sample

```
['that', 'picture', 'i', 'just', 'seen', 'whoa', 'dere', '!', '!',  
 '>', '>', '>', '>', '>', '>', '>']
```


Exercise 4 - count_words

You won't use all the tokens (words) appearing in the data for training. Instead, you will use the more frequently used words.

- You will focus on the words that appear at least N times in the data

Hints

```
In [15]: # UNQ_C4 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
### GRADED_FUNCTION: count_words ###
def count_words(tokenized_sentences):
    """
    Count the number of word appearance in the tokenized sentences

    Args:
        tokenized_sentences: List of lists of strings

    Returns:
        dict that maps word (str) to the frequency (int)
    """

    word_counts = {}
    ### START CODE HERE (Replace instances of 'None' with your code) ###

    # Loop through each sentence
    for sentence in tokenized_sentences: # complete this line

        # Go through each token in the sentence
        for token in sentence: # complete this line

            # If the token is not in the dictionary yet, set the count to 1
            if token not in word_counts: # complete this line
                word_counts[token] = 1

            # If the token is already in the dictionary, increment the count by 1
            else:
                word_counts[token] += 1

    ### END CODE HERE ###

    return word_counts
```

```
In [16]: # test your code
tokenized_sentences = [['sky', 'is', 'blue', '.'],
                       ['leaves', 'are', 'green', '.'],
                       ['roses', 'are', 'red', '.']]
count_words(tokenized_sentences)
```

```
Out[16]: {'sky': 1,
          'is': 1,
          'blue': 1,
          '.': 3,
          'leaves': 1,
          'are': 2,
          'green': 1,
          'roses': 1,
          'red': 1}
```

Expected output

Note that the order may differ.

```
{'sky': 1,
 'is': 1,
 'blue': 1,
 '.': 3,
 'leaves': 1,
 'are': 2,
 'green': 1,
 'roses': 1,
 'red': 1}
```

```
In [17]: # Test your function
w3_unittest.test_count_words(count_words)
```

All tests passed

Handling 'Out of Vocabulary' words

If your model is performing autocomplete, but encounters a word that it never saw during training, it won't have an input word to help it determine the next word to suggest. The model will not be able to predict the next word because there are no counts for the current word.

- This 'new' word is called an 'unknown word', or **out of vocabulary (OOV)** words.
- The percentage of unknown words in the test set is called the **OOV** rate.

To handle unknown words during prediction, use a special token to represent all unknown words 'unk'.

- Modify the training data so that it has some 'unknown' words to train on.
- Words to convert into "unknown" words are those that do not occur very frequently in the training set.
- Create a list of the most frequent words in the training set, called the **closed vocabulary**.
- Convert all the other words that are not part of the closed vocabulary to the token 'unk'.

Exercise 5 - get_words_with_nplus_frequency

You will now create a function that takes in a text document and a threshold `count_threshold`.

- Any word whose count is greater than or equal to the threshold `count_threshold` is kept in the closed vocabulary.
- Returns the word closed vocabulary list.

```
In [18]: # UNQ_C5 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
### GRADED_FUNCTION: get_words_with_nplus_frequency ###
def get_words_with_nplus_frequency(tokenized_sentences, count_threshold):
    """
    Find the words that appear N times or more

    Args:
        tokenized_sentences: List of lists of sentences
        count_threshold: minimum number of occurrences for a word to be in the

    Returns:
        List of words that appear N times or more
    """
    # Initialize an empty list to contain the words that
    # appear at least 'minimum_freq' times.
    closed_vocab = []

    # Get the word counts of the tokenized sentences
    # Use the function that you defined earlier to count the words
    word_counts = count_words(tokenized_sentences)

    ### START CODE HERE (Replace instances of 'None' with your code) ###

    # for each word and its count
    for word, cnt in word_counts.items(): # complete this line

        # check that the word's count
        # is at least as great as the minimum count
        if cnt >= count_threshold:

            # append the word to the list
            closed_vocab.append(word)
    ### END CODE HERE ###

    return closed_vocab
```

```
In [19]: # test your code
tokenized_sentences = [['sky', 'is', 'blue', '.'],
                       ['leaves', 'are', 'green', '.'],
                       ['roses', 'are', 'red', '.']]
tmp_closed_vocab = get_words_with_nplus_frequency(tokenized_sentences, count_t
print(f"Closed vocabulary:")
print(tmp_closed_vocab)
```

Closed vocabulary:
['.', 'are']

Expected output

Closed vocabulary:
['.', 'are']

```
In [20]: # Test your function
w3_unittest.test_get_words_with_nplus_frequency(get_words_with_nplus_frequency
```

All tests passed

Exercise 6 - replace_oov_words_by_unk

The words that appear `count_threshold` times or more are in the closed vocabulary.

- All other words are regarded as `unknown` .
- Replace words not in the closed vocabulary with the token `<unk>` .

```
In [21]: # UNQ_C6 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
### GRADED_FUNCTION: replace_oov_words_by_unk ###
def replace_oov_words_by_unk(tokenized_sentences, vocabulary, unknown_token="<unk>"):
    """
    Replace words not in the given vocabulary with the unknown token.

    Args:
        tokenized_sentences: List of lists of strings
        vocabulary: List of strings that we will use
        unknown_token: A string representing unknown (out-of-vocabulary) words

    Returns:
        List of lists of strings, with words not in the vocabulary replaced
    """

    # Use a set for faster lookup
    vocabulary = set(vocabulary)

    replaced_tokenized_sentences = []

    for sentence in tokenized_sentences:
        replaced_sentence = []
        for token in sentence:
            if token in vocabulary:
                replaced_sentence.append(token)
            else:
                replaced_sentence.append(unknown_token) # ✓ Use the parameter
        replaced_tokenized_sentences.append(replaced_sentence)

    return replaced_tokenized_sentences
```

```
In [22]: tokenized_sentences = [["dogs", "run"], ["cats", "sleep"]]
vocabulary = ["dogs", "sleep"]
tmp_replaced_tokenized_sentences = replace_oov_words_by_unk(tokenized_sentences, vocabulary, unknown_token="<unk>")
print(f"Original sentence:")
print(tokenized_sentences)
print(f"tokenized_sentences with less frequent words converted to '<unk>':")
print(tmp_replaced_tokenized_sentences)
```

Original sentence:

```
[['dogs', 'run'], ['cats', 'sleep']]
```

tokenized_sentences with less frequent words converted to '<unk>':

```
[['dogs', '<unk>'], ['<unk>', 'sleep']]
```

Expected answer

Original sentence:

```
[['dogs', 'run'], ['cats', 'sleep']]
```

tokenized_sentences with less frequent words converted to '<unk>':

```
[['dogs', '<unk>'], ['<unk>', 'sleep']]
```

```
In [23]: # Test your function
w3_unittest.test_replace_oov_words_by_unk(replace_oov_words_by_unk)
```

All tests passed

Exercise 7 - preprocess_data

Now we are ready to process our data by combining the functions that you just implemented.

1. Find tokens that appear at least `count_threshold` times in the training data.
2. Replace tokens that appear less than `count_threshold` times by "<unk>" both for training and test data.

```
In [24]: # UNQ_C7 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
### GRADED_FUNCTION: preprocess_data ###
def preprocess_data(train_data, test_data, count_threshold, unknown_token="<unk>")
    """
    Preprocess data:
        - Find tokens that appear at least count_threshold times in train_data
        - Replace low-frequency tokens with unknown_token in both train and test data

    Returns:
        - preprocessed training data
        - preprocessed test data
        - vocabulary list
    """
    # Step 1: Get closed vocabulary
    vocabulary = get_words_with_nplus_frequency(train_data, count_threshold)

    # Step 2: Replace rare words with unknown_token in train and test data
    train_data_replaced = replace_oov_words_by_unk(train_data, vocabulary, unknown_token)
    test_data_replaced = replace_oov_words_by_unk(test_data, vocabulary, unknown_token)

    return train_data_replaced, test_data_replaced, vocabulary
```

```
In [25]: # test your code
tmp_train = [['sky', 'is', 'blue', '.'],
             ['leaves', 'are', 'green']]
tmp_test = [['roses', 'are', 'red', '.']]

tmp_train_repl, tmp_test_repl, tmp_vocab = preprocess_data(tmp_train,
                                                         tmp_test,
                                                         count_threshold = 1
                                                         )

print("tmp_train_repl")
print(tmp_train_repl)
print()
print("tmp_test_repl")
print(tmp_test_repl)
print()
print("tmp_vocab")
print(tmp_vocab)
```

```
tmp_train_repl
[['sky', 'is', 'blue', '.'], ['leaves', 'are', 'green']]
```

```
tmp_test_repl
[['<unk>', 'are', '<unk>', '.']]
```

```
tmp_vocab
['sky', 'is', 'blue', '.', 'leaves', 'are', 'green']
```

Expected outcome

```
tmp_train_repl
[['sky', 'is', 'blue', '.'], ['leaves', 'are', 'green']]
```

```
tmp_test_repl
[['<unk>', 'are', '<unk>', '.']]
```

```
tmp_vocab
['sky', 'is', 'blue', '.', 'leaves', 'are', 'green']
```

```
In [26]: # Test your function
w3_unittest.test_preprocess_data(preprocess_data)
```

All tests passed

Preprocess the train and test data

Run the cell below to complete the preprocessing both for training and test sets.

```
In [27]: minimum_freq = 2
train_data_processed, test_data_processed, vocabulary = preprocess_data(train_data,
                                                                    test_data,
                                                                    minimum_freq)
```

```
In [28]: print("First preprocessed training sample:")
print(train_data_processed[0])
print()
print("First preprocessed test sample:")
print(test_data_processed[0])
print()
print("First 10 vocabulary:")
print(vocabulary[0:10])
print()
print("Size of vocabulary:", len(vocabulary))
```

First preprocessed training sample:

```
['i', 'personally', 'would', 'like', 'as', 'our', 'official', 'glove', 'of',
'the', 'team', 'local', 'company', 'and', 'quality', 'production']
```

First preprocessed test sample:

```
['that', 'picture', 'i', 'just', 'seen', 'whoa', 'dere', '!', '!', '>', '>',
'>', '>', '>', '>', '>']
```

First 10 vocabulary:

```
['i', 'personally', 'would', 'like', 'as', 'our', 'official', 'glove', 'of',
'the']
```

Size of vocabulary: 14821

Expected output

First preprocessed training sample:

```
['i', 'personally', 'would', 'like', 'as', 'our', 'official', 'glov
e', 'of', 'the', 'team', 'local', 'company', 'and', 'quality', 'produ
ction']
```

First preprocessed test sample:

```
['that', 'picture', 'i', 'just', 'seen', 'whoa', 'dere', '!', '!',
'>', '>', '>', '>', '>', '>', '>']
```

First 10 vocabulary:

```
['i', 'personally', 'would', 'like', 'as', 'our', 'official', 'glov
e', 'of', 'the']
```

Size of vocabulary: 14821

You are done with the preprocessing section of the assignment. Objects

`train_data_processed` , `test_data_processed` , and `vocabulary` will be used in the rest of the exercises.

2 - Develop n-gram based Language Models

In this section, you will develop the n-grams language model.

- Assume the probability of the next word depends only on the previous n-gram.

Processing math: 100%

The previous n-gram is the series of the previous 'n' words.

The conditional probability for the word at position 't' in the sentence, given that the words preceding it are $w_{t-n} \dots w_{t-2}, w_{t-1}$ is:

$$P(w_t | w_{t-n} \dots w_{t-1})$$

You can estimate this probability by counting the occurrences of these series of words in the training data.

- The probability can be estimated as a ratio, where
- The numerator is the number of times word 't' appears after words t-n through t-1 appear in the training data.
- The denominator is the number of times word t-n through t-1 appears in the training data.

$$\hat{P}(w_t | w_{t-n} \dots w_{t-1}) = \frac{C(w_{t-n} \dots w_{t-1}, w_t)}{C(w_{t-n} \dots w_{t-1})}$$

- The function $C(\dots)$ denotes the number of occurrence of the given sequence.
- \hat{P} means the estimation of P .
- Notice that denominator of the equation (2) is the number of occurrence of the previous n words, and the numerator is the same sequence followed by the word w_t .

Later, you will modify the equation (2) by adding k-smoothing, which avoids errors when any counts are zero.

The equation (2) tells us that to estimate probabilities based on n-grams, you need the counts of n-grams (for denominator) and (n+1)-grams (for numerator).

Exercise 8 - count_n_grams

Next, you will implement a function that computes the counts of n-grams for an arbitrary number n .

When computing the counts for n-grams, prepare the sentence beforehand by prepending $n - 1$ starting markers "<s>" to indicate the beginning of the sentence.

- For example, in the tri-gram model ($n=3$), a sequence with two start tokens "<s>" should predict the first word of a sentence.
- So, if the sentence is "I like food", modify it to be "<s> <s> I like food".
- Also prepare the sentence for counting by appending an end token "<e>" so that the model can predict when to finish a sentence.

Technical note: In this implementation, you will store the counts as a dictionary.

- The key of each key-value pair in the dictionary is a **tuple** of n words (and not a list)
- The value in the key-value pair is the number of occurrences.
- The reason for using a tuple as a key instead of a list is because a list in Python is a mutable object (it can be changed after it is first created). A tuple is "immutable", so it cannot be altered after it is first created. This makes a tuple suitable as a data type for the key in a dictionary.
- Although for a n-gram you need to use $n-1$ starting markers for a sentence, you will want to prepend n starting markers in order to use them to compute the initial probability for the $(n+1)$ -gram later in the assignment.

Hints

```
In [29]: # UNQ_C8 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
### GRADED FUNCTION: count_n_grams ###
def count_n_grams(data, n, start_token='<s>', end_token = '<e>'):
    """
    Count all n-grams in the data

    Args:
        data: List of lists of words
        n: number of words in a sequence

    Returns:
        A dictionary that maps a tuple of n-words to its frequency
    """

    # Initialize dictionary of n-grams and their counts
    n_grams = {}

    ### START CODE HERE (ReplacE instances of 'None' with your code) ###

    # Go through each sentence in the data
    for sentence in data: # complete this line

        # prepend start token n times, and append <e> one time
        sentence = [start_token] * n + sentence + [end_token]

        # convert list to tuple
        # So that the sequence of words can be used as
        # a key in the dictionary
        sentence = tuple(sentence)

        # Use 'i' to indicate the start of the n-gram
        # from index 0
        # to the last index where the end of the n-gram
        # is within the sentence.

        for i in range(len(sentence)-n+1): # complete this line

            # Get the n-gram from i to i+n
            n_gram = sentence[i:i+n]

            # check if the n-gram is in the dictionary
            if n_gram in n_grams: # complete this line

                # Increment the count for this n-gram
                n_grams[n_gram] += 1
            else:
                # Initialize this n-gram count to 1
                n_grams[n_gram] = 1

    ### END CODE HERE ###
    return n_grams
```

```
In [30]: # test your code
# CODE REVIEW COMMENT: Outcome does not match expected outcome
sentences = [['i', 'like', 'a', 'cat'],
              ['this', 'dog', 'is', 'like', 'a', 'cat']]
print("Uni-gram:")
print(count_n_grams(sentences, 1))
print("Bi-gram:")
print(count_n_grams(sentences, 2))
```

```
Uni-gram:
{(' <s> ',): 2, ('i',): 1, ('like',): 2, ('a',): 2, ('cat',): 2, (' <e> ',): 2,
 ('this',): 1, ('dog',): 1, ('is',): 1}
Bi-gram:
{(' <s> ', ' <s> '): 2, (' <s> ', 'i'): 1, ('i', 'like'): 1, ('like', 'a'): 2,
 ('a', 'cat'): 2, ('cat', ' <e> '): 2, (' <s> ', 'this'): 1, ('this', 'dog'): 1,
 ('dog', 'is'): 1, ('is', 'like'): 1}
```

Expected outcome:

```
Uni-gram:
{(' <s> ',): 2, ('i',): 1, ('like',): 2, ('a',): 2, ('cat',): 2, (' <e> ',): 2,
 ('this',): 1, ('dog',): 1, ('is',): 1}
Bi-gram:
{(' <s> ', ' <s> '): 2, (' <s> ', 'i'): 1, ('i', 'like'): 1, ('like', 'a'): 2,
 ('a', 'cat'): 2, ('cat', ' <e> '): 2, (' <s> ', 'this'): 1, ('this', 'dog'): 1,
 ('dog', 'is'): 1, ('is', 'like'): 1}
```

Take a look to the (' <s> ', ' <s> ') element in the bi-gram dictionary. Although for a bi-gram you will only require one starting mark, as in the element (' <s> ', 'i'), this (' <s> ', ' <s> ') element will be helpful when computing the probabilities using tri-grams (the corresponding count will be used as denominator).

```
In [31]: # Test your function
w3_unittest.test_count_n_grams(count_n_grams)
```

All tests passed

Exercise 9 - estimate_probability

Next, estimate the probability of a word given the prior 'n' words using the n-gram counts.

$$\hat{P}(w_t | w_{t-n} \dots w_{t-1}) = \frac{C(w_{t-n} \dots w_{t-1}, w_t)}{C(w_{t-n} \dots w_{t-1})}$$

This formula doesn't work when a count of an n-gram is zero..

- Suppose we encounter an n-gram that did not occur in the training data.
- Then, the equation (2) cannot be evaluated (it becomes zero divided by zero).

A way to handle zero counts is to add k-smoothing.

- K-smoothing adds a positive constant k to each numerator and $k \times |V|$ in the denominator, where $|V|$ is the number of words in the vocabulary.

$$\hat{P}(w_t | w_{t-n} \dots w_{t-1}) = \frac{C(w_{t-n} \dots w_{t-1}, w_t) + k}{C(w_{t-n} \dots w_{t-1}) + k |V|}$$

For n-grams that have a zero count, the equation (3) becomes $\frac{1}{|V|}$.

- This means that any n-gram with zero count has the same probability of $\frac{1}{|V|}$.

Define a function that computes the probability estimate (3) from n-gram counts and a constant k .

- The function takes in a dictionary 'n_gram_counts', where the key is the n-gram and the value is the count of that n-gram.
- The function also takes another dictionary n_plus1_gram_counts, which you'll use to find the count for the previous n-gram plus the current word.

Hints

```

In [48]: # UNQ_C9 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
        ### GRADED FUNCTION: estimate_probability ###
        def estimate_probability(word, previous_n_gram,
                                n_gram_counts, n_plus1_gram_counts, vocabulary_size,
                                """
            Estimate the probabilities of a next word using the n-gram counts with k-s

            Args:
                word: next word
                previous_n_gram: A sequence of words of length n
                n_gram_counts: Dictionary of counts of n-grams
                n_plus1_gram_counts: Dictionary of counts of (n+1)-grams
                vocabulary_size: number of words in the vocabulary
                k: positive constant, smoothing parameter

            Returns:
                A probability
            """
            # convert list to tuple to use it as a dictionary key
            previous_n_gram = tuple(previous_n_gram)

            ### START CODE HERE (Replace instances of 'None' with your code) ###

            # Set the denominator
            # If the previous n-gram exists in the dictionary of n-gram counts,
            # Get its count. Otherwise set the count to zero
            # Use the dictionary that has counts for n-grams
            previous_n_gram_count = n_gram_counts.get(previous_n_gram, 0)

            # Calculate the denominator using the count of the previous n gram
            # and apply k-smoothing
            denominator = previous_n_gram_count + (k * vocabulary_size)

            # Define n plus 1 gram as the previous n-gram plus the current word as a t
            n_plus1_gram = previous_n_gram + (word, )

            # Set the count to the count in the dictionary,
            # otherwise 0 if not in the dictionary
            # use the dictionary that has counts for the n-gram plus current word
            n_plus1_gram_count = n_plus1_gram_counts.get(n_plus1_gram, 0)

            # Define the numerator use the count of the n-gram plus current word,
            # and apply smoothing
            numerator = n_plus1_gram_count + k

            # Calculate the probability as the numerator divided by denominator
            probability = numerator / denominator

            ### END CODE HERE ###

            return probability

```

```
In [49]: # test your code
sentences = [['i', 'like', 'a', 'cat'],
              ['this', 'dog', 'is', 'like', 'a', 'cat']]
unique_words = list(set(sentences[0] + sentences[1]))

unigram_counts = count_n_grams(sentences, 1)
bigram_counts = count_n_grams(sentences, 2)
tmp_prob = estimate_probability("cat", ["a"], unigram_counts, bigram_counts, 1)

print(f"The estimated probability of word 'cat' given the previous n-gram 'a' is: 0.3333")
```

The estimated probability of word 'cat' given the previous n-gram 'a' is: 0.3333

Expected output

The estimated probability of word 'cat' given the previous n-gram 'a' is: 0.3333

```
In [50]: # Test your function
w3_unittest.test_estimate_probability(estimate_probability)

All tests passed
```

Estimate probabilities for all words

The function defined below loops over all words in vocabulary to calculate probabilities for all possible words.

- This function is provided for you.

```
In [51]: def estimate_probabilities(previous_n_gram, n_gram_counts, n_plus1_gram_counts,
    """
    Estimate the probabilities of next words using the n-gram counts with k-smoothing.

    Args:
        previous_n_gram: A sequence of words of length n
        n_gram_counts: Dictionary of counts of n-grams
        n_plus1_gram_counts: Dictionary of counts of (n+1)-grams
        vocabulary: List of words
        k: positive constant, smoothing parameter

    Returns:
        A dictionary mapping from next words to the probability.
    """
    # convert list to tuple to use it as a dictionary key
    previous_n_gram = tuple(previous_n_gram)

    # add <e> <unk> to the vocabulary
    # <s> is not needed since it should not appear as the next word
    vocabulary = vocabulary + [end_token, unknown_token]
    vocabulary_size = len(vocabulary)

    probabilities = {}
    for word in vocabulary:
        probability = estimate_probability(word, previous_n_gram,
                                         n_gram_counts, n_plus1_gram_counts,
                                         vocabulary_size, k=k)

        probabilities[word] = probability

    return probabilities
```

```
In [52]: # test your code
sentences = [['i', 'like', 'a', 'cat'],
             ['this', 'dog', 'is', 'like', 'a', 'cat']]
unique_words = list(set(sentences[0] + sentences[1]))
unigram_counts = count_n_grams(sentences, 1)
bigram_counts = count_n_grams(sentences, 2)

estimate_probabilities(["a"], unigram_counts, bigram_counts, unique_words, k=1)
```

```
Out[52]: {'this': 0.09090909090909091,
          'like': 0.09090909090909091,
          'is': 0.09090909090909091,
          'dog': 0.09090909090909091,
          'cat': 0.2727272727272727,
          'i': 0.09090909090909091,
          'a': 0.09090909090909091,
          '<e>': 0.09090909090909091,
          '<unk>': 0.09090909090909091}
```

Expected output

```
{'cat': 0.2727272727272727,
 'i': 0.09090909090909091,
 'this': 0.09090909090909091,
 'a': 0.09090909090909091,
 'is': 0.09090909090909091,
 'like': 0.09090909090909091}
```

```
In [53]: # Additional test
trigram_counts = count_n_grams(sentences, 3)
estimate_probabilities(["<s>", "<s>"], bigram_counts, trigram_counts, unique_v
```

```
Out[53]: {'this': 0.18181818181818182,
 'like': 0.09090909090909091,
 'is': 0.09090909090909091,
 'dog': 0.09090909090909091,
 'cat': 0.09090909090909091,
 'i': 0.18181818181818182,
 'a': 0.09090909090909091,
 '<e>': 0.09090909090909091,
 '<unk>': 0.09090909090909091}
```

Expected output

```
{'cat': 0.09090909090909091,
 'i': 0.18181818181818182,
 'this': 0.18181818181818182,
 'a': 0.09090909090909091,
 'is': 0.09090909090909091,
 'like': 0.09090909090909091,
 'dog': 0.09090909090909091,
 '<e>': 0.09090909090909091,
 '<unk>': 0.09090909090909091}
```

Count and probability matrices

As we have seen so far, the n-gram counts computed above are sufficient for computing the probabilities of the next word.

- It can be more intuitive to present them as count or probability matrices.
- The functions defined in the next cells return count or probability matrices.
- This function is provided for you.

```
In [54]: def make_count_matrix(n_plus1_gram_counts, vocabulary):
# add <e> <unk> to the vocabulary
# <s> is omitted since it should not appear as the next word
vocabulary = vocabulary + ["<e>", "<unk>"]

# obtain unique n-grams
n_grams = []
for n_plus1_gram in n_plus1_gram_counts.keys():
    n_gram = n_plus1_gram[0:-1]
    n_grams.append(n_gram)
n_grams = list(set(n_grams))

# mapping from n-gram to row
row_index = {n_gram:i for i, n_gram in enumerate(n_grams)}
# mapping from next word to column
col_index = {word:j for j, word in enumerate(vocabulary)}

nrow = len(n_grams)
ncol = len(vocabulary)
count_matrix = np.zeros((nrow, ncol))
for n_plus1_gram, count in n_plus1_gram_counts.items():
    n_gram = n_plus1_gram[0:-1]
    word = n_plus1_gram[-1]
    if word not in vocabulary:
        continue
    i = row_index[n_gram]
    j = col_index[word]
    count_matrix[i, j] = count

count_matrix = pd.DataFrame(count_matrix, index=n_grams, columns=vocabulary)
return count_matrix
```

```
In [55]: sentences = [['i', 'like', 'a', 'cat'],
                     ['this', 'dog', 'is', 'like', 'a', 'cat']]
unique_words = list(set(sentences[0] + sentences[1]))
bigram_counts = count_n_grams(sentences, 2)

print('bigram counts')
display(make_count_matrix(bigram_counts, unique_words))
```

bigram counts

	this	like	is	dog	cat	i	a	<e>	<unk>
(a,)	0.0	0.0	0.0	0.0	2.0	0.0	0.0	0.0	0.0
(this,)	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
(like,)	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0	0.0
(<s>,)	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
(is,)	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
(cat,)	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0
(dog,)	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
(i,)	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

bigram counts

	cat	i	this	a	is	like	dog	<e>	<unk>
(<s>,)	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
(a,)	2.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
(this,)	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
(like,)	0.0	0.0	0.0	2.0	0.0	0.0	0.0	0.0	0.0
(dog,)	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
(cat,)	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0
(is,)	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
(i,)	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0

```
In [56]: # Show trigram counts
print('\ntrigram counts')
trigram_counts = count_n_grams(sentences, 3)
display(make_count_matrix(trigram_counts, unique_words))
```

trigram counts

	this	like	is	dog	cat	i	a	<e>	<unk>
(is, like)	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
(i, like)	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0
(like, a)	0.0	0.0	0.0	0.0	2.0	0.0	0.0	0.0	0.0
(this, dog)	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
(<s>, this)	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
(<s>, i)	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
(dog, is)	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
(a, cat)	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0
(<s>, <s>)	1.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0

Expected output

trigram counts

	cat	i	this	a	is	like	dog	<e>	<unk>
(dog, is)	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
(this, dog)	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0
(a, cat)	0.0	0.0	0.0	0.0	0.0	0.0	0.0	2.0	0.0
(like, a)	2.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
(is, like)	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
(<s>, i)	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0
(i, like)	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0
(<s>, <s>)	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
(<s>, this)	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0

The following function calculates the probabilities of each word given the previous n-gram, and stores this in matrix form.

• This function is provided for you.

```
In [57]: def make_probability_matrix(n_plus1_gram_counts, vocabulary, k):
count_matrix = make_count_matrix(n_plus1_gram_counts, unique_words)
count_matrix += k
prob_matrix = count_matrix.div(count_matrix.sum(axis=1), axis=0)
return prob_matrix
```

```
In [58]: sentences = [['i', 'like', 'a', 'cat'],
['this', 'dog', 'is', 'like', 'a', 'cat']]
unique_words = list(set(sentences[0] + sentences[1]))
bigram_counts = count_n_grams(sentences, 2)
print("bigram probabilities")
display(make_probability_matrix(bigram_counts, unique_words, k=1))
```

bigram probabilities

	this	like	is	dog	cat	i	a	<e>	<unk>
(a,)	0.090909	0.090909	0.090909	0.090909	0.272727	0.090909	0.090909	0.090909	0.090909
(this,)	0.100000	0.100000	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000
(like,)	0.090909	0.090909	0.090909	0.090909	0.090909	0.090909	0.272727	0.090909	0.090909
(<s>,)	0.181818	0.090909	0.090909	0.090909	0.090909	0.181818	0.090909	0.090909	0.090909
(is,)	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000
(cat,)	0.090909	0.090909	0.090909	0.090909	0.090909	0.090909	0.090909	0.272727	0.090909
(dog,)	0.100000	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000
(i,)	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000



```
In [59]: print("trigram probabilities")
trigram_counts = count_n_grams(sentences, 3)
display(make_probability_matrix(trigram_counts, unique_words, k=1))
```

trigram probabilities

	this	like	is	dog	cat	i	a	<e>	<unk>
(is, like)	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000	0.200000	0.100000	0.100000
(i, like)	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000	0.200000	0.100000	0.100000
(like, a)	0.090909	0.090909	0.090909	0.090909	0.272727	0.090909	0.090909	0.090909	0.090909
(this, dog)	0.100000	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000
(<s>, this)	0.100000	0.100000	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000
(<s>, i)	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000
(dog, is)	0.100000	0.200000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000	0.100000
(a, cat)	0.090909	0.090909	0.090909	0.090909	0.090909	0.090909	0.090909	0.272727	0.090909
(<s>, <s>)	0.181818	0.090909	0.090909	0.090909	0.090909	0.181818	0.090909	0.090909	0.090909

Confirm that you obtain the same results as for the `estimate_probabilities` function that you implemented.

3 - Perplexity

In this section, you will generate the perplexity score to evaluate your model on the test set.

- You will also use back-off when needed.
- Perplexity is used as an evaluation metric of your language model.
- To calculate the perplexity score of the test set on an n-gram model, use:

$$PP(W) = \sqrt[N]{\prod_{t=n+1}^N \frac{1}{P(w_t | w_{t-n} \dots w_{t-1})}}$$

- where N is the length of the sentence.
- n is the number of words in the n-gram (e.g. 2 for a bigram).
- In math, the numbering starts at one and not zero.

In code, array indexing starts at zero, so the code will use ranges for t according to this formula:

$$PP(W) = \sqrt[N]{\prod_{t=n}^{N-1} \frac{1}{P(w_t | w_{t-n} \dots w_{t-1})}}$$

Exercise 10 - calculate_perplexity

Compute the perplexity score given an N-gram count matrix and a sentence.

Note: For the sake of simplicity, in the code below, <s> is included in perplexity score calculation.

Hints

```
In [60]: def calculate_perplexity(sentence, n_gram_counts, n_plus1_gram_counts,
                                vocabulary_size, k=1.0, start_token="<s>", end_token="<e>"):
    """
    Calculate perplexity for a sentence using Laplace smoothing.

    Args:
        sentence: List of strings (tokens)
        n_gram_counts: Dictionary of counts of n-grams
        n_plus1_gram_counts: Dictionary of counts of (n+1)-grams
        vocabulary_size: Number of unique words in the vocabulary
        k: Smoothing constant (default = 1.0)
        start_token: Token that marks the start of a sentence (default = "<s>")
        end_token: Token that marks the end of a sentence (default = "<e>")

    Returns:
        Perplexity score (float)
    """
    if not n_gram_counts:
        raise ValueError("n_gram_counts dictionary is empty.")

    n = len(next(iter(n_gram_counts))) # get n from the first n-gram key
    sentence = [start_token] * n + sentence + [end_token] # pad sentence
    N = len(sentence)

    sum_log_prob = 0.0

    for t in range(n, N):
        n_gram = tuple(sentence[t - n:t])
        word = sentence[t]
        n_gram_count = n_gram_counts.get(n_gram, 0)
        n_plus1_gram = n_gram + (word,)
        n_plus1_gram_count = n_plus1_gram_counts.get(n_plus1_gram, 0)

        probability = (n_plus1_gram_count + k) / (n_gram_count + k * vocabulary_size)
        sum_log_prob += -math.log(probability)

    perplexity = math.exp(sum_log_prob / N)
    return perplexity
```

In [61]: *# test your code*

```
sentences = [['i', 'like', 'a', 'cat'],
              ['this', 'dog', 'is', 'like', 'a', 'cat']]
unique_words = list(set(sentences[0] + sentences[1]))

unigram_counts = count_n_grams(sentences, 1)
bigram_counts = count_n_grams(sentences, 2)

perplexity_train = calculate_perplexity(sentences[0],
                                       unigram_counts, bigram_counts,
                                       len(unique_words), k=1.0)
print(f"Perplexity for first train sample: {perplexity_train:.4f}")

test_sentence = ['i', 'like', 'a', 'dog']
perplexity_test = calculate_perplexity(test_sentence,
                                       unigram_counts, bigram_counts,
                                       len(unique_words), k=1.0)
print(f"Perplexity for test sample: {perplexity_test:.4f}")
```

Perplexity for first train sample: 2.8040
Perplexity for test sample: 3.9654

In [62]: *# Test your function*

```
w3_unittest.test_calculate_perplexity(calculate_perplexity)
```

All tests passed

Expected Output

Perplexity for first train sample: 2.8040
Perplexity for test sample: 3.9654

Note: If your sentence is really long, there will be underflow when multiplying many fractions.

- To handle longer sentences, modify your implementation to take the sum of the log of the probabilities.

4 - Build an Auto-complete System

In this section, you will combine the language models developed so far to implement an auto-complete system.

Exercise 11 - suggest_a_word

Compute probabilities for all possible next words and suggest the most likely one.

- This function also take an optional argument `start_with`, which specifies the first few letters of the next words.

Hints


```

In [63]: # UNQ_C11 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
# GRADED FUNCTION: suggest_a_word
def suggest_a_word(previous_tokens, n_gram_counts, n_plus1_gram_counts, vocabulary):
    """
    Get suggestion for the next word

    Args:
        previous_tokens: The sentence you input where each token is a word. Must be a list.
        n_gram_counts: Dictionary of counts of (n+1)-grams
        n_plus1_gram_counts: Dictionary of counts of (n+1)-grams
        vocabulary: List of words
        k: positive constant, smoothing parameter
        start_with: If not None, specifies the first few letters of the next word

    Returns:
        A tuple of
        - string of the most likely next word
        - corresponding probability
    """

    # Length of previous words
    n = len(list(n_gram_counts.keys())[0])

    # From the words that the user already typed
    # get the most recent 'n' words as the previous n-gram
    previous_n_gram = previous_tokens[-n:]

    # Estimate the probabilities that each word in the vocabulary
    # is the next word,
    # given the previous n-gram, the dictionary of n-gram counts,
    # the dictionary of n plus 1 gram counts, and the smoothing constant
    probabilities = estimate_probabilities(previous_n_gram,
                                          n_gram_counts, n_plus1_gram_counts,
                                          vocabulary, k=k)

    # Initialize suggested word to None
    # This will be set to the word with highest probability
    suggestion = None

    # Initialize the highest word probability to 0
    # this will be set to the highest probability
    # of all words to be suggested
    max_prob = 0

    ### START CODE HERE (Replace instances of 'None' with your code) ###

    # For each word and its probability in the probabilities dictionary:
    for word, prob in probabilities.items(): # complete this line

        # If the optional start_with string is set
        if start_with is not None: # complete this line

            # Check if the beginning of word does not match with the letters in start_with
            if not word.startswith(start_with): # complete this line

                # if they don't match, skip this word (move onto the next word)
                continue # complete this line

            # Check if this word's probability
            # is greater than the current maximum probability
            if prob > max_prob: # complete this line

```



```

        # If so, save this word as the best suggestion (so far)
        suggestion = word

        # Save the new maximum probability
        max_prob = prob

    ### END CODE HERE

    return suggestion, max_prob

```

```

In [64]: # test your code
sentences = [['i', 'like', 'a', 'cat'],
              ['this', 'dog', 'is', 'like', 'a', 'cat']]
unique_words = list(set(sentences[0] + sentences[1]))

unigram_counts = count_n_grams(sentences, 1)
bigram_counts = count_n_grams(sentences, 2)

previous_tokens = ["i", "like"]
tmp_suggest1 = suggest_a_word(previous_tokens, unigram_counts, bigram_counts,
print(f"The previous words are 'i like',\n\tand the suggested word is `{tmp_su

print()
# test your code when setting the starts_with
tmp_starts_with = 'c'
tmp_suggest2 = suggest_a_word(previous_tokens, unigram_counts, bigram_counts,
print(f"The previous words are 'i like', the suggestion must start with `{tmp_

```

The previous words are 'i like',
and the suggested word is `a` with a probability of 0.2727

The previous words are 'i like', the suggestion must start with `c`
and the suggested word is `cat` with a probability of 0.0909

Expected output

The previous words are 'i like',
and the suggested word is `a` with a probability of 0.2727

The previous words are 'i like', the suggestion must start with `c`
and the suggested word is `cat` with a probability of 0.0909

```

In [65]: # Test your function
w3_unittest.test_suggest_a_word(suggest_a_word)

```

All tests passed

Get multiple suggestions

The function defined below loops over various n-gram models to get multiple suggestions.

```
In [66]: def get_suggestions(previous_tokens, n_gram_counts_list, vocabulary, k=1.0, start_with=None):
    model_counts = len(n_gram_counts_list)
    suggestions = []
    for i in range(model_counts-1):
        n_gram_counts = n_gram_counts_list[i]
        n_plus1_gram_counts = n_gram_counts_list[i+1]

        suggestion = suggest_a_word(previous_tokens, n_gram_counts,
                                    n_plus1_gram_counts, vocabulary,
                                    k=k, start_with=start_with)
        suggestions.append(suggestion)
    return suggestions
```

```
In [67]: # test your code
sentences = [['i', 'like', 'a', 'cat'],
             ['this', 'dog', 'is', 'like', 'a', 'cat']]
unique_words = list(set(sentences[0] + sentences[1]))

unigram_counts = count_n_grams(sentences, 1)
bigram_counts = count_n_grams(sentences, 2)
trigram_counts = count_n_grams(sentences, 3)
quadgram_counts = count_n_grams(sentences, 4)
qintgram_counts = count_n_grams(sentences, 5)

n_gram_counts_list = [unigram_counts, bigram_counts, trigram_counts, quadgram_counts, qintgram_counts]
previous_tokens = ["i", "like"]
tmp_suggest3 = get_suggestions(previous_tokens, n_gram_counts_list, unique_words)

print(f"The previous words are 'i like', the suggestions are:")
display(tmp_suggest3)
```

The previous words are 'i like', the suggestions are:

```
[('a', 0.2727272727272727),
 ('a', 0.2),
 ('this', 0.1111111111111111),
 ('this', 0.1111111111111111)]
```

Suggest multiple words using n-grams of varying length

Congratulations! You have developed all building blocks for implementing your own auto-complete systems.

Let's see this with n-grams of varying lengths (unigrams, bigrams, trigrams, 4-grams...6-grams).

```
In [68]: n_gram_counts_list = []
for n in range(1, 6):
    print("Computing n-gram counts with n =", n, "...")
    n_model_counts = count_n_grams(train_data_processed, n)
    n_gram_counts_list.append(n_model_counts)
```

```
Computing n-gram counts with n = 1 ...
Computing n-gram counts with n = 2 ...
Computing n-gram counts with n = 3 ...
Computing n-gram counts with n = 4 ...
Computing n-gram counts with n = 5 ...
```

```
In [69]: previous_tokens = ["i", "am", "to"]
tmp_suggest4 = get_suggestions(previous_tokens, n_gram_counts_list, vocabulary)

print(f"The previous words are {previous_tokens}, the suggestions are:")
display(tmp_suggest4)
```

The previous words are ['i', 'am', 'to'], the suggestions are:

```
[('be', 0.027665685098338604),
 ('have', 0.00013487086115044844),
 ('have', 0.00013490725126475548),
 ('i', 6.746272684341901e-05)]
```

```
In [70]: previous_tokens = ["i", "want", "to", "go"]
tmp_suggest5 = get_suggestions(previous_tokens, n_gram_counts_list, vocabulary)

print(f"The previous words are {previous_tokens}, the suggestions are:")
display(tmp_suggest5)
```

The previous words are ['i', 'want', 'to', 'go'], the suggestions are:

```
[('to', 0.014051961029228078),
 ('to', 0.004697942168993581),
 ('to', 0.0009424436216762033),
 ('to', 0.0004044489383215369)]
```

```
In [71]: previous_tokens = ["hey", "how", "are"]
tmp_suggest6 = get_suggestions(previous_tokens, n_gram_counts_list, vocabulary)

print(f"The previous words are {previous_tokens}, the suggestions are:")
display(tmp_suggest6)
```

The previous words are ['hey', 'how', 'are'], the suggestions are:

```
[('you', 0.023426812585499317),
 ('you', 0.003559435862995299),
 ('you', 0.00013491635186184566),
 ('i', 6.746272684341901e-05)]
```

```
In [72]: previous_tokens = ["hey", "how", "are", "you"]
tmp_suggest7 = get_suggestions(previous_tokens, n_gram_counts_list, vocabulary)

print(f"The previous words are {previous_tokens}, the suggestions are:")
display(tmp_suggest7)
```

The previous words are ['hey', 'how', 'are', 'you'], the suggestions are:

```
[('re', 0.023973994311255586),
 ('?', 0.002888465830762161),
 ('?', 0.0016134453781512605),
 ('<e>', 0.00013491635186184566)]
```

```
In [73]: previous_tokens = ["hey", "how", "are", "you"]
tmp_suggest8 = get_suggestions(previous_tokens, n_gram_counts_list, vocabulary)

print(f"The previous words are {previous_tokens}, the suggestions are:")
display(tmp_suggest8)
```

The previous words are ['hey', 'how', 'are', 'you'], the suggestions are:

```
[('do', 0.009020723283218204),
 ('doing', 0.0016411737674785006),
 ('doing', 0.00047058823529411766),
 ('dvd', 6.745817593092283e-05)]
```

Congratulations!



You've completed this assignment by building an autocomplete model using an n-gram language model!

Please continue onto the fourth and final week of this course!

IT PC x 19A! x Histo x Auto x Natu x Part x Dow x c2-w x c2-w x s1.p x Dow x s2.p x s1.p x




← → ↺ coursera.org/learn/probabilistic-models-in-nlp/programming/7GrBr/autocomplete

☆ ↻ ⬇ ⚙

Search in course

Search

English   M 

Hide menu

Lecture: Autocomplete

Lecture Notes (Optional)

Practice Quiz

Assignment: Autocomplete

✓

 Programming Assignment: Autocomplete 3h

Natural Language Processing with Probabilistic Models > Week 3 > Autocomplete

< Previous Next >

Programming Assignment: Autocomplete

✓

 Passed • 110/110 points

Deadline Pass this assignment by May 4, 11:59 PM PDT

Launch Notebook

Instructions

 My submissions

Complete the assignment within the Jupyter environment and Submit via the 'Submit' button on the top-right of the page.


Like

Dislike

Report an issue

88°F Mostly clear

Search



ENG IN

21:05 30-04-2025